

Transaction

Transaction: Collection of operations that form a single logical unit of work is called transaction.

A transaction is a unit of program execution that accesses and possibly updates various data items.

Transaction access data using two operations:

read(x): transfers data item X from the database to a logical buffer belonging to the transaction that executed the read operation.

write(x): transfers the data item X from the local buffer of the transaction that executed the write back to the database.

Eg; Transfer Rs. 500/- from account A to account B

```
T1
read(A);
A = A - 500;
write(A);
read(B)
B = B + 500;
write(B);
```

ACID Properties

solution

- A - Atomicity :
- C - Consistency
- I - Isolation
- D - Durability

Atomicity :- Either all operations of the transaction are reflected properly in the database, or none are.

T₁

read (A)
A = A - 500
write (A)
read (B)
B = B + 500
write (B)

Initial value of A is Rs. 2000/-
Initial value of B is Rs. 3000/-

failures

generates inconsistent state
A = 1500
B = 3000

Consistency :- Execution of transaction preserves the consistency of the database if the operations are executed in isolation. (with no other transaction executing concurrently)

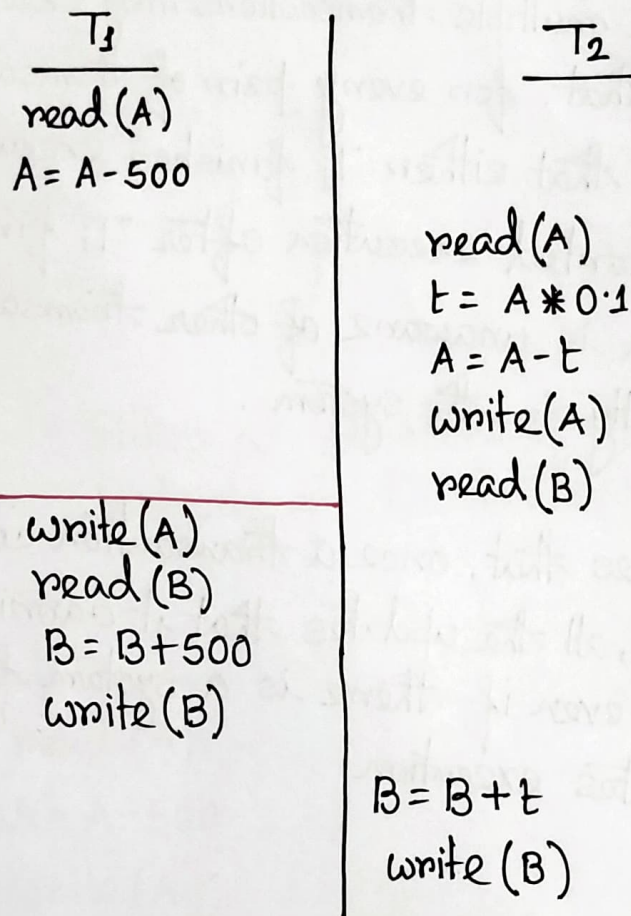
sum of A and B (before transaction) = sum of A and B (after transaction)

Isolation: Perform every unit of operations isolately.

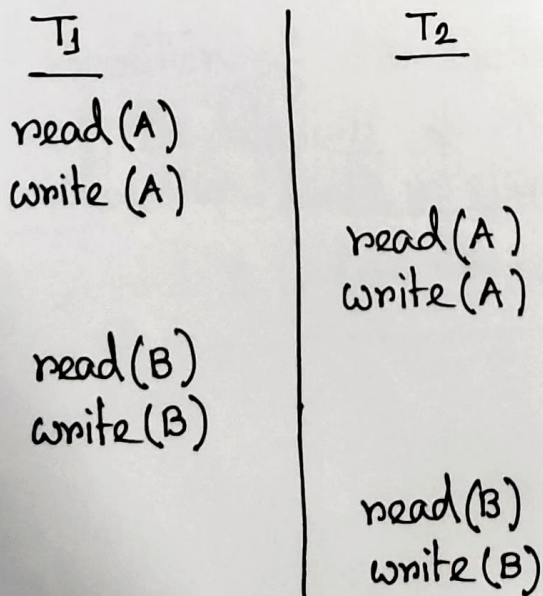
Even though multiple transactions may execute concurrently the system guarantees that, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished. Thus, each transaction is unaware of other transactions executing concurrently in the system.

Durability: It guarantees that, once a transaction completes successfully, all the updates that it carried out on the database persist, even if there is a system failure after the transaction completes execution.

Serializability



Any problem in schedule S ??



Concurrent Executions

T_1 : Transfer Rs. 500/- from account A to account B

T_2 : Transfer 10% of account A to account B

Serial Schedule (Serial Executions)

S_1

<u>T_1</u>	<u>T_2</u>
read(A)	
$A = A - 500$	
write(A)	
read(B)	
$B = B + 500$	
write(B)	
	read(A)
	$t = A * 0.1$
	$A = A - t$
	write(A)
	read(B)
	$B = B + t$
	write(B)

S_2

<u>T_1</u>	<u>T_2</u>
	read(A)
	$t = A * 0.1$
	$A = A - t$
	write(A)
	read(B)
	$B = B + t$
	write(B)
read(A)	
$A = A - 500$	
write(A)	
read(B)	
$B = B + 500$	
write(B)	

concurrent schedule (Concurrent Execution) :

serial


S_1'		S_2'	
<u>T_1</u>	<u>T_2</u>	<u>T_1</u>	<u>T_2</u>
read(A) $A = A - 500$ write(A)	read(A) $t = A * 0.1$ $A = A - t$ write(A)	read(A) $A = A - 500$ write(A)	read(A) $t = A * 0.1$ $A = A - t$ write(A)
read(B) $B = B + 500$ write(B)	read(B) $B = B + t$ write(B)	read(B) $B = B + 500$ write(B)	read(B) $B = B + t$ write(B)

→ Improved throughput and Resource utilization

→ Reduced waiting time.

Serial Schedule: A schedule S is said to be serial if for every participating transaction T in the schedule, executes consecutively.

Serializable Schedule: A schedule S of n transactions is serializable, if it is equivalent to some serial schedule of the same n transactions.



Conflict Serializability

Consider a schedule S in which there are two consecutive instructions I_i and I_j , of transactions T_i and T_j respectively.

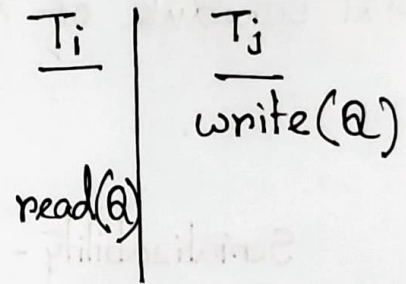

Three Cases:

1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. The order of I_i and I_j does not matter, since the same value of Q read by T_i and T_j , regardless of the order

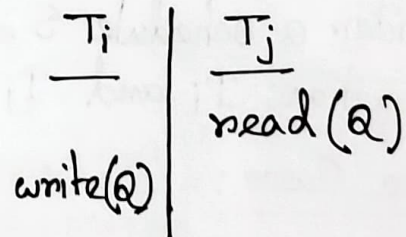
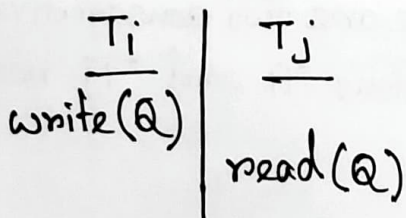
<u>T_i</u>	<u>T_j</u>
read(Q)	read(Q)

<u>T_i</u>	<u>T_j</u>
read(Q)	read(Q)



2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. If I_i comes before I_j ,
 then T_i does not read the value of Q that is written by
 T_j in instruction I_j . If I_j comes before I_i then
 T_i reads the value of Q that is written by T_j .
 Thus the order of I_i and I_j matters



3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. The order of I_i and I_j matters
 for the same reason as the previous case.



4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. In this case the value
 obtained by the next read(Q) instruction of S is affected,
 since the result of only the latter of the two write instructions
 is preserved in the database.



* Instructions I_i and I_j conflict if they are operations by different transactions on the same data item, and at least one of these instructions is a write operation.

*** If instructions I_i and I_j refer to different data items, then they do not conflict and we can swap I_i and I_j without affecting the results of any instructions in the schedule.

Conflict Equivalent: If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent.

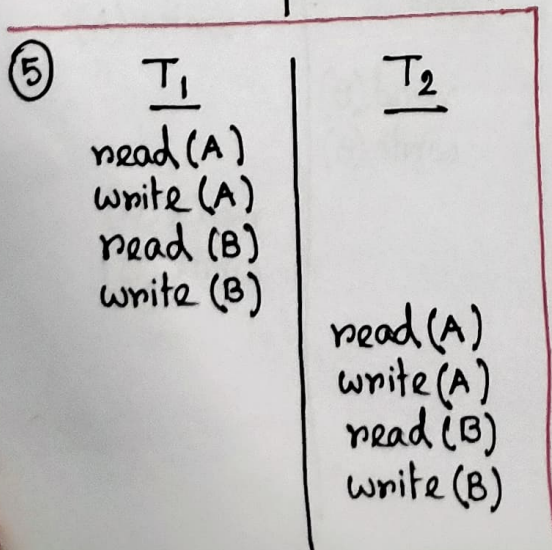
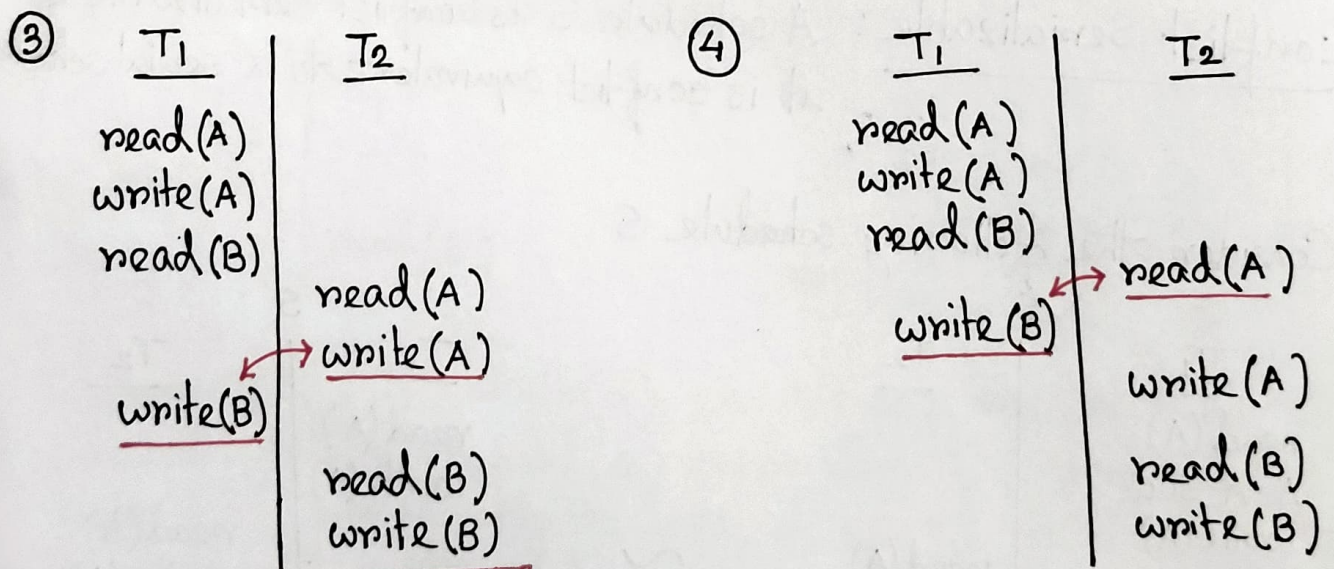
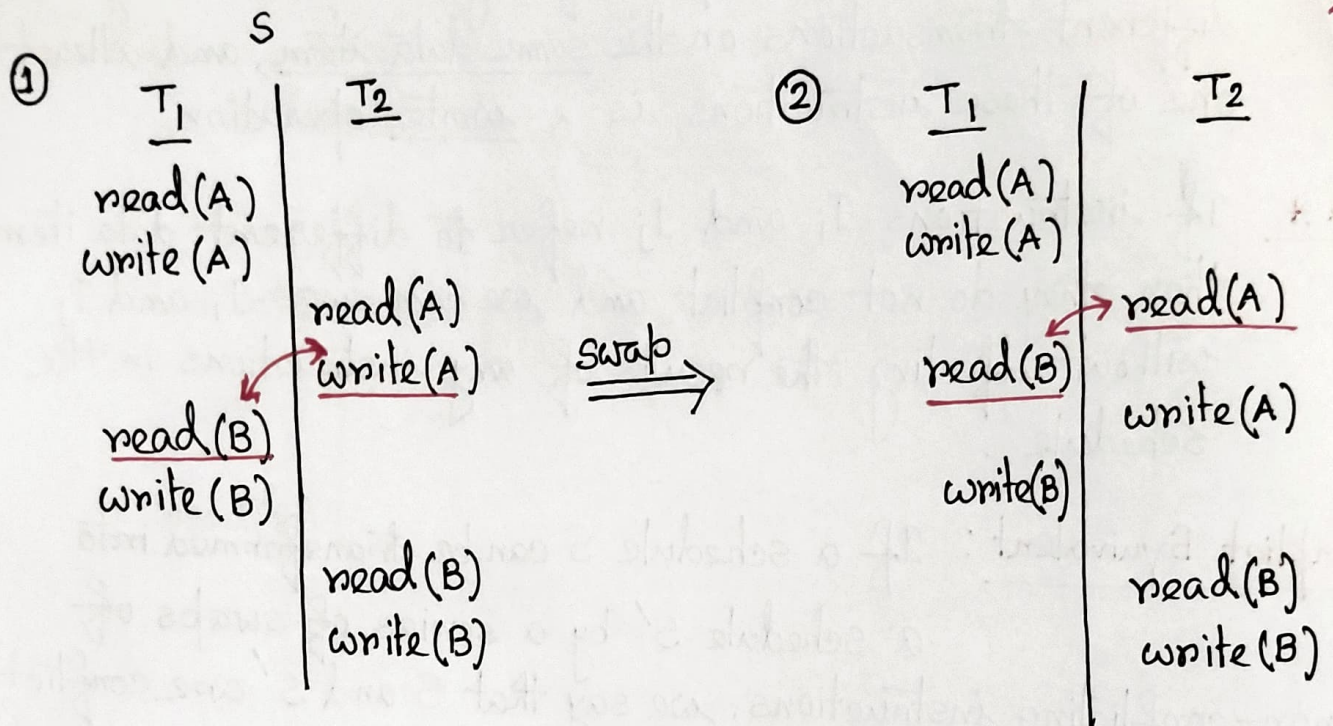
Conflict Serializable: A schedule S is conflict serializable if it is conflict equivalent to a serial schedule.

Consider the following schedule S .

S			S	
<u>T_1</u>	<u>T_2</u>		<u>T_1</u>	<u>T_2</u>
read(A) $A = A - 500$ write(A)			read(A) write(A)	
	read(A) $t = A * 0.1$ $A = A - t$ write(A)	\approx		read(A) write(A)
read(B) $B = B + 500$ write(B)			read(B) write(B)	
	read(B) $B = B + t$ write(B)			read(B) write(B)

* Check the schedule S is conflict serializable or not.

view Ser



\Rightarrow Serial Schedule, so schedule S is a conflict serializable schedule.

View Serializability: A schedule S is view serializable if it is view equivalent to a serial schedule.

View Equivalent: Consider two schedules S and S' , where the same set of transactions participates in both schedules. The schedules S and S' are said to be view equivalent if three conditions are met:

1. For each data item Q , if transaction T_i reads the initial value of Q in schedule S , then transaction T_i must, in schedule S' , also read the initial value of Q .
2. For each data item Q , if transaction T_i executes $\text{read}(Q)$ in schedule S , and if that value was produced by a $\text{write}(Q)$ operation executed by transaction T_j , then the $\text{read}(Q)$ operation of transaction T_i must, in schedule S' , also read the value of Q that was produced by the same $\text{write}(Q)$ operation of transaction T_j .
3. For each data item Q , the transaction (if any) that performs the final $\text{write}(Q)$ operation in schedule S must perform the final $\text{write}(Q)$ operation in schedule S' .

* Consider the following schedules S and S' . Check whether S and S' are view equivalent or not. Check whether S' is view serializable or not.

S (Serial)

<u>T_1</u>	<u>T_2</u>	<u>T_3</u>
read(Q)		
write(Q)		
	write(Q)	
		write(Q)

S' (Concurrent)

<u>T_1</u>	<u>T_2</u>	<u>T_3</u>
read(Q)		
write(Q)	write(Q)	
		write(Q)


Testing for Serializability

Precedence graph is used to test the serializability of a schedule.

A precedence graph consists of a pair $G = (V, E)$, where V is the set of vertices and E is the set of Edges. The set of vertices consists of all the transactions participating in the schedule.

The set of edges consists of all edges $T_i \rightarrow T_j$ for which one of the three conditions holds:

1. T_i executes $write(Q)$ before T_j executes $read(Q)$.
2. T_i executes $read(Q)$ before T_j executes $write(Q)$.
3. T_i executes $write(Q)$ before T_j executes $write(Q)$.



**** If the precedence graph contains no cycles, then the schedule S is conflict serializable, otherwise schedule S is not conflict serializable**

*** Consider the following schedules and test for serializability**

<u>S_1</u>		<u>S_2</u>	
<u>T_1</u>	<u>T_2</u>	<u>T_1</u>	<u>T_2</u>
$r(A)$ $w(A)$	$r(A)$ $w(A)$	$r(A)$	$r(A)$ $w(A)$ $r(B)$
$r(B)$ $w(B)$	$r(B)$ $w(B)$	$w(A)$ $r(B)$ $w(B)$	$w(B)$

241

S₁



S₂

